(a) Hidden-state conflict

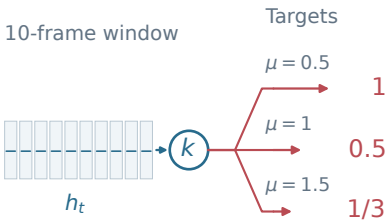

At this key, hidden gain changes the teacher's target.

Illustration: $a_x = 0.5/\mu$

(b) Probe and certify

PRIVILEGED REPLAY

source reset → decision → brake

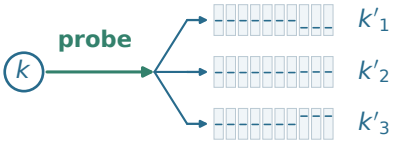

Probe separates visible inputs

Cross paths

Y_{rm}

No universal candidate

Sequential mean fails

Probe continuation succeeds

Exact binding · gain · consistency

(c) Evidence-based routing

SOURCE-ONLY ROUTING

loss

Certified probe → **INFO** ✓

Consistent / safe → **EXEC** ✓

Conflict / invalid → **ESC** ×

Keep INFO + EXEC targets

Drop targets; keep input history

TRAIN Retained source targets only

DEPLOY $h_t \longrightarrow \pi_\theta(h_t) \longrightarrow a_t$

Contract schematic

H = 10 · fresh state per decision · no teacher or replay input